

Sapienza University of Rome

Master in Artificial Intelligence and Robotics

Machine Learning

A.Y. 2025/2026

Prof. Luca Iocchi

7. Linear models for classification

Luca Iocchi

Overview

- Linearly separable data
- Linear models
- Least squares
- Perceptron
- Fisher's linear discriminant
- Support Vector Machines

References

C. Bishop. Pattern Recognition and Machine Learning. Sect. 4.1, 7.1

T. Mitchell. Machine Learning. Section 4.4

Linear Models for Classification


Learning a function $f : X \rightarrow Y$, with ...

- $X \subseteq \Re^d$
- $Y = \{C_1, \dots, C_k\}$

assuming *linearly separable* data.

Linearly separable data

Instances in a data set are *linearly separable* iff there exists a hyperplane that separates the instance space into two regions, such that differently classified instances are separated



Linear discriminant functions

Linear discriminant function

$$y : X \rightarrow \{C_1, \dots, C_K\}$$

Two classes:

$$y(x) = w^T x + w_0$$

K -classes:

$$y_1(x) = w_1^T x + w_{10}$$

...

$$y_K(x) = w_K^T x + w_{K0}$$

Compact notation

Two classes:

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}, \text{ with:}$$


$$\tilde{\mathbf{w}} = \begin{pmatrix} w_0 \\ \mathbf{w} \end{pmatrix}, \tilde{\mathbf{x}} = \begin{pmatrix} 1 \\ \mathbf{x} \end{pmatrix}$$

K classes:

$$y(\mathbf{x}) = \begin{pmatrix} y_1(\mathbf{x}) \\ \dots \\ y_K(\mathbf{x}) \end{pmatrix} = \begin{pmatrix} \mathbf{w}_1^T \mathbf{x} + w_{10} \\ \dots \\ \mathbf{w}_K^T \mathbf{x} + w_{K0} \end{pmatrix} = \begin{pmatrix} \tilde{\mathbf{w}}_1^T \\ \dots \\ \tilde{\mathbf{w}}_K^T \end{pmatrix} \tilde{\mathbf{x}} = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}, \text{ with:}$$

$$\tilde{\mathbf{W}}^T = \begin{pmatrix} \tilde{\mathbf{w}}_1^T \\ \dots \\ \tilde{\mathbf{w}}_K^T \end{pmatrix}, \text{ i.e.: } \tilde{\mathbf{W}} = (\tilde{\mathbf{w}}_1, \dots, \tilde{\mathbf{w}}_K)$$

Linear discriminant functions



Multiple classes

K -class discriminant comprising K linear functions

$$y(x) = \begin{pmatrix} y_1(x) \\ \vdots \\ y_K(x) \end{pmatrix} = \begin{pmatrix} \tilde{w}_1^T \tilde{x} \\ \vdots \\ \tilde{w}_K^T \tilde{x} \end{pmatrix} = \tilde{W}^T \tilde{x}$$


Classify x as C_k with $k = \operatorname{argmax}_{j=1,\dots,K} \{y_j(x)\}$

Decision boundary between C_k and C_j (hyperplane in \Re^{D-1}):

$$(\tilde{w}_k - \tilde{w}_j)^T \tilde{x} = 0$$

Multiple classes

Example of K -class discriminant



Learning linear discriminants

Given a multi-class classification problem and data set D with linearly separable data,

determine \tilde{W} such that $y(x) = \tilde{W}^T \tilde{x}$ is the K-class discriminant.

Approaches to learn linear discriminants

- Least squares
- Perceptron
- Fisher's linear discriminant
- Support Vector Machines

Least squares

Given $D = \{(x_n, t_n)_{n=1}^N\}$, find the linear discriminant

$$y(x) = \tilde{W}^T \tilde{x}$$

1-of-K coding scheme for t : $x \in C_k \rightarrow t_k = 1, t_j = 0$ for all $j \neq k$.
 E.g., $t_n = (0, \dots, 1, \dots, 0)^T$

$$\tilde{x} = \begin{pmatrix} \tilde{x}_1^T \\ \vdots \\ \tilde{x}_N^T \end{pmatrix} \quad T = \begin{pmatrix} t_1^T \\ \vdots \\ t_N^T \end{pmatrix}$$

Least squares

Minimize sum-of-squares error function (removed $\tilde{\cdot}$)

$$E(W) = \frac{1}{2} \operatorname{Tr} \left\{ (XW - T)^T (XW - T) \right\}$$

Closed-form solution:

$$W^* = \underbrace{(X^T X)^{-1} X^T}_{X^\dagger} T$$

Learned model:

$$y(X) = W^{*T} X = T^T (X^\dagger)^T X$$

Least squares

Classification of new instance $x' \notin D$

Use learnt W^* to compute:



$$y(x') = W^{*T} x' = \begin{pmatrix} y_1(x') \\ \dots \\ y_K(x') \end{pmatrix}$$

Predict C_k with:

$$k = \operatorname{argmax}_{j \in \{1, \dots, K\}} \{y_j(x')\}$$


Issues with least squares

Assume Gaussian conditional distributions. Not robust to outliers!



Perceptron

Learning $f : \mathbb{R}^d \rightarrow \{-1, +1\}$



$$o(x_1, \dots, x_d) = \begin{cases} 1 & \text{if } w_0 + w_1 x_1 + \dots + w_d x_d > 0 \\ -1 & \text{otherwise.} \end{cases}$$

$$o(x) = \begin{cases} 1 & \text{if } w^T x > 0 \\ -1 & \text{otherwise.} \end{cases} = sign(w^T x)$$

Perceptron training rule

Consider the *unthresholded linear unit*, where

$$o = w_0 + w_1 x_1 + \dots + w_d x_d = w^T x$$

Let's learn w_i from training examples $D = \{(x_n, t_n)\}_{n=1}^N\}$ that minimize the squared error (*loss function*)

$$E(w) \equiv \frac{1}{2} \sum_{n=1}^N (t_n - o_n)^2 = \frac{1}{2} \sum_{n=1}^N (t_n - w^T x_n)^2$$

Perceptron training rule

$$\begin{aligned}
 \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{n=1}^N (t_n - w_n^T x_n)^2 = \frac{1}{2} \sum_{n=1}^N \frac{\partial}{\partial w_i} (t_n - w^T x_n)^2 \\
 &= \frac{1}{2} \sum_{n=1}^N 2(t_n - w^T x_n) \frac{\partial}{\partial w_i} (t_n - w^T x_n) \\
 &= \sum_{n=1}^N (t_n - w^T x_n) \frac{\partial}{\partial w_i} (t_n - w^T x_n) \\
 &= \sum_{n=1}^N (t_n - w^T x_n)(-x_{i,n})
 \end{aligned}$$

Perceptron training rule

Unthresholded unit:

Update of weights w

$$\begin{aligned}
 w_i &\leftarrow w_i + \Delta w_i \\
 \Delta w_i &= -\eta \frac{\partial E}{\partial w_i} = \eta \sum_{n=1}^N (t_n - w^T x_n) x_{i,n}
 \end{aligned}$$

η is a small constant (e.g., 0.05) called *learning rate*

Perceptron training rule

Thresholded unit:

Update of weights w

$$w_i \leftarrow w_i + \Delta w_i$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i} = \eta \sum_{n=1}^N (t_n - \text{sign}(w^T x_n)) x_{i,n}$$

Perceptron algorithm

Given perceptron model $o(x) = \text{sign}(w^T x)$ and data set D , determine weights w.

- ① Initialize \hat{w} (e.g. small random values)
- ② Repeat until termination condition
 - $\hat{w}_i \leftarrow \hat{w}_i + \Delta w_i$
- ③ Output \hat{w}

Perceptron algorithm

Batch mode: Consider all dataset D

$$\Delta w_i = \eta \sum_{(x,t) \in D} (t - o(x)) x_i$$

Mini-Batch mode: Choose a small subset $S \subset D$

$$\Delta w_i = \eta \sum_{(x,t) \in S} (t - o(x)) x_i$$

Incremental mode: Choose one sample $(x, t) \in D$

$$\Delta w_i = \eta (t - o(x)) x_i$$

$o(x) = w^T x$ for unthresholded, $o(x) = sign(w^T x)$ for thresholded
 Incremental and mini-batch modes speed up convergence and are less sensitive to local minima.

Perceptron algorithm

Termination conditions

- Predefined number of iterations
- Threshold on changes in the loss function $E(w)$


Perceptron training rule

Example:

$$\eta = 0.1, x_i = 0.8$$

- if $t = 1$ and $o = -1$ then $\Delta w_i = 0.16$
- if $t = -1$ and $o = 1$ then $\Delta w_i = -0.16$

Perceptron training rule



Perceptron training rule

Can prove it will converge:

- if training data is linearly separable
- and η sufficiently small

Small $\eta \rightarrow$ slow convergence.

Perceptron: Prediction

Classification of new instance $x' \notin D$:

Predict C_k , with $k = \text{sign}(\hat{w}^T x')$, using learnt \hat{w}

Fisher's linear discriminant

Adjusting w to find a direction that maximizes class separation.


Consider a data set with N_1 points in C_1 and N_2 points in C_2

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in C_1} \mathbf{x}_n \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in C_2} \mathbf{x}_n$$

Choose w that maximizes $J(w) = w^T(\mathbf{m}_2 - \mathbf{m}_1)$, subject to $\|w\| = 1$.

Fisher's linear discriminant

$$w \propto (\mathbf{m}_2 - \mathbf{m}_1)$$



Fisher's linear discriminant

Fisher criterion

$$J(w) = \frac{w^T S_B w}{w^T S_W w}$$

with

$$S_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

Between class scatter

$$S_W = \sum_{n \in C_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in C_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

Within class scatter

Choose w that maximizes $J(w)$.

Fisher's linear discriminant

Find w that maximizes

$$J(w) = \frac{w^T S_B w}{w^T S_W w}$$


by solving

$$\frac{d}{dw} J(w) = 0$$

$$\Rightarrow w^* \propto S_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

Fisher's linear discriminant

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$



Fisher's linear discriminant

Summarizing, given a two classes classification problem, Fisher's linear discriminant is given by the function $y = \mathbf{w}^T \mathbf{x}$ and the classification of new instances is given by $y \geq -w_0$
where

$$\mathbf{w} = \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

$$w_0 = \mathbf{w}^T \mathbf{m}$$

\mathbf{m} is the global mean of all the data set.

Fisher's linear discriminant

Multiple classes.

$$y = W^T x$$


Maximizing

$$J(W) = \text{Tr} \left\{ (\mathbf{W} S_W \mathbf{W}^T)^{-1} (\mathbf{W} S_B \mathbf{W}^T) \right\}$$

...

Support Vector Machines

Support Vector Machines (SVM) for Classification aims at maximum margin providing for better accuracy.



Support Vector Machines

Let's consider binary classification $f : X \rightarrow \{+1, -1\}$ with data set $D = \{(x_n, t_n)_{n=1}^N\}$, $t_n \in \{+1, -1\}$ and a linear model

$$y(x) = w^T x + w_0$$

Assume D is linearly separable

$$\exists w, w_0 \text{ s.t. } \begin{cases} y(x_n) > 0, & \text{if } t_n = +1 \\ y(x_n) < 0, & \text{if } t_n = -1 \end{cases}$$

$$t_n y(x_n) > 0 \quad \forall n = 1, \dots, N$$

Support Vector Machines

Let x_k be the closest point of the data set D to the hyperplane $h : w^T x + w_0 = 0$

the *margin* (smallest distance between x_k and h) is $\frac{|y(x_k)|}{\|w\|}$

Given data set D and hyperplane h , the margin is computed as

$$\min_{n=1, \dots, N} \frac{|y(x_n)|}{\|w\|} = \dots = \frac{1}{\|w\|} \min_{n=1, \dots, N} [t_n (w^T x_n + w_0)]$$

using the property $|y(x_n)| = t_n y(x_n)$

Support Vector Machines

Given data set D , the hyperplane $h^* : \mathbf{w}^{*T} \mathbf{x} + w_0^* = 0$ with maximum margin is computed as

$$\mathbf{w}^*, w_0^* = \operatorname{argmax}_{\mathbf{w}, w_0} \frac{1}{\|\mathbf{w}\|} \min_{n=1,\dots,N} [t_n(\mathbf{w}^T \mathbf{x}_n + w_0)]$$

Support Vector Machines

Rescaling all the points does not affect the solution.

Rescale in such a way that for the closest point \mathbf{x}_k we have

$$t_k(\mathbf{w}^T \mathbf{x}_k + w_0) = 1$$

Canonical representation:


$$t_n(\mathbf{w}^T \mathbf{x}_n + w_0) \geq 1 \quad \forall n = 1, \dots, N$$

Support Vector Machines

When the maximum margin hyperplane w^*, w_0^* is found, there will be at least 2 closest points x_k^\oplus and x_k^\ominus (one for each class).

$$w^* T x_k^\oplus + w_0^* = +1$$

$$w^* T x_k^\ominus + w_0^* = -1$$



Support Vector Machines

In the canonical representation of the problem the maximum margin hyperplane can be found by solving the optimization problem

$$w^*, w_0^* = \operatorname{argmax}_{\|w\|} \frac{1}{\|w\|} = \operatorname{argmin} \frac{1}{2} \|w\|^2$$

subject to

$$t_n(w^T x_n + w_0) \geq 1 \quad \forall n = 1, \dots, N$$

Quadratic programming problem solved with Lagrangian method.

Support Vector Machines

Solution

$$\mathbf{w}^* = \sum_{n=1}^N a_n^* t_n \mathbf{x}_n$$

a_i^* (Lagrange multipliers): results of the Lagrangian optimization problem

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m$$

subject to

$$a_n \geq 0 \quad \forall n = 1, \dots, N$$

$$\sum_{n=1}^N a_n t_n = 0$$

Support Vector Machines

Note:

samples \mathbf{x}_n for which $a_n^* = 0$ do not contribute to the solution

Karush-Kuhn-Tucker (KKT) condition:

for each $\mathbf{x}_n \in D$, either $a_n^* = 0$ or $t_n y(\mathbf{x}_n) = 1$

thus $t_n y(\mathbf{x}_n) > 1$ implies $a_n^* = 0$

Support vectors: \mathbf{x}_k such that $t_k y(\mathbf{x}_k) = 1$ and $a_k^* > 0$

$$SV \equiv \{\mathbf{x}_k \in D \mid t_k y(\mathbf{x}_k) = 1\}$$

Support Vector Machines

Hyperplanes expressed with support vectors

$$y(x) = \sum_{x_j \in SV} a_j^* t_j x_j^T x + w_0^* = 0$$

Note: other vectors $x_n \notin SV$ do not contribute ($a_n^* = 0$)

Support Vector Machines

To compute w_0^* :

Support vector $x_k \in SV$ satisfies $t_k y(x_k) = 1$

$$t_k \left(\sum_{x_j \in SV} a_j^* t_j x_k^T x_j + w_0^* \right) = 1$$

Multiplying by t_k and using $t_k^2 = 1$

$$w_0^* = t_k - \sum_{x_j \in SV} a_j^* t_j x_k^T x_j$$

Support Vector Machines

Instead of using one particular support vector x_k to determine w_0

$$w_0^* = t_k - \sum_{x_j \in SV} a_j^* t_j x_k^T x_j$$

a more stable solution is obtained by averaging over all the support vectors

$$w_0^* = \frac{1}{|SV|} \sum_{x_k \in SV} \left(t_k - \sum_{x_j \in S} a_j^* t_j x_k^T x_j \right)$$

Support Vector Machines

Given the maximum margin hyperplane determined by a_k^* , w_0^*

Classification of a new instance x'

$$y(x') = sign \left(\sum_{x_k \in SV} a_k^* t_k x'^T x_k + w_0^* \right)$$

Support Vector Machines

Optimization problem for determining w, w_0 (dimension $d + 1$, with $X = \mathbb{R}^d$) transformed in an optimization problem for determining a (dimension $|D|$)

Efficient when $d \ll |D|$ (most of a_i will be zero).
Very useful when d is large or infinite.


Support Vector Machines with soft margin constraints

What if data are “almost” linearly separable (e.g., a few points are on the “wrong side”)

Let us introduce *slack variables* $\xi_n \geq 0$ $n = 1, \dots, N$

Support Vector Machines with soft margin constraints

- $\xi_n = 0$ if point on or inside the correct margin boundary
- $0 < \xi_n \leq 1$ if point inside the margin but correct side
- $\xi_n > 1$ if point on wrong side of boundary



when $\xi_n = 1$, the sample lies on the decision boundary $y(x_n) = 0$
 when $\xi_n > 1$, the sample will be mis-classified

Support Vector Machines with soft margin constraints

Soft margin constraint

$$t_n y(x_n) \geq 1 - \xi_n, \quad n = 1, \dots, N$$

Optimization problem with soft margin constraints

$$\mathbf{w}^*, w_0^* = \operatorname{argmin} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

subject to

$$t_n y(x_n) \geq 1 - \xi_n, \quad n = 1, \dots, N$$

$$\xi_n \geq 0, \quad n = 1, \dots, N$$

C is a constant (inverse of a regularization coefficient)

Support Vector Machines with soft margin constraints

Solution similar to the case of linearly separable data.

$$\mathbf{w}^* = \sum_{n=1}^N a_n^* t_n \mathbf{x}_n$$

$$w_0^* = \dots$$

with a_n^* computed as solution of a Lagrangian optimization problem.

Basis functions



So far we considered models working directly on \mathbf{x} .

All the results hold if we consider a non-linear transformation of the inputs $\phi(\mathbf{x})$ (*basis functions*).

Decision boundaries will be linear in the feature space ϕ and non-linear in the original space \mathbf{x}

Classes that are linearly separable in the feature space ϕ may not be separable in the input space \mathbf{x} .

Basis functions example



Basis functions examples

- Linear
- Polynomial
- Radial Basis Function (RBF)
- Sigmoid
- ...

Linear models for non-linear functions

Learning non-linear function

$$f : X \rightarrow \{C_1, \dots, C_K\}$$

from data set D non-linearly separable.

Find a non-linear transformation ϕ and learn a linear model

$$y(x) = w^T \phi(x) + w_0 \text{ (*two classes*)}$$

$$y_k(x) = w_k^T \phi(x) + w_{k0} \text{ (*multiple classes*)}$$

Summary

- Basic methods for learning linear classification functions
- Based on solution of an optimization problem
- Closed form vs. iterative solutions
- Sensitivity to outliers
- Learning non-linear functions with linear models using basis functions
- Further developed as kernel methods